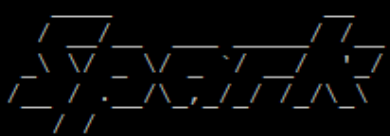


## Практическое задание №6

Запускаем spark, добавляем пакет cassandra-connector

```
student782_4@bigdataanalytics2-worker-shdpt-v31-1-0:~  
login as: student782_4  
Authenticating with public key "imported-openssh-key"  
Last login: Thu Aug 19 09:05:15 2021 from mail.softwarecom.ru  
[student782_4@bigdataanalytics2-head-shdpt-v31-1-0 ~]$ ssh 10.0.0.19  
Last login: Thu Aug 19 09:05:24 2021 from bigdataanalytics2-head-shdpt-v31-1-0.n  
ovalocal  
[student782_4@bigdataanalytics2-worker-shdpt-v31-1-0 ~]$ /spark2.4/bin/pyspark -  
-packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5,com.datastax.spark:sp  
ark-cassandra-connector_2.11:2.4.2  
Python 2.7.5 (default, Nov 16 2020, 22:23:17)  
[GCC 4.8.5 20150623 (Red Hat 4.8.5-44)] on linux2  
Type "help", "copyright", "credits" or "license" for more information.  
Warning: Ignoring non-Spark config property: hive.metastore.uris  
Ivy Default Cache set to: /home/student782_4/.ivy2/cache  
The jars for the packages stored in: /home/student782_4/.ivy2/jars  
:: loading settings :: url = jar:file:/spark2.4/jars/ivy-2.4.0.jar!/org/apache/i  
vy/core/settings/ivysettings.xml  
org.apache.spark#spark-sql-kafka-0-10_2.11 added as a dependency  
com.datastax.spark#spark-cassandra-connector_2.11 added as a dependency  
:: resolving dependencies :: org.apache.spark#spark-submit-parent-f8e7cec4-837e-  
4b06-90b9-d0a82704face;1.0  
   confs: [default]  
   found org.apache.spark#spark-sql-kafka-0-10_2.11;2.4.5 in central  
   found org.apache.kafka#kafka-clients;2.0.0 in central
```

Создадим конфигурацию spark

```
Welcome to  
 version 2.4.7  
  
Using Python version 2.7.5 (default, Nov 16 2020 22:23:17)  
SparkSession available as 'spark'.  
>>> from pyspark.sql import SparkSession  
>>> from pyspark.sql import functions as F  
>>> from pyspark.sql.types import StringType, StructType  
>>>  
>>> spark.conf.set("spark.cassandra.connection.host", "localhost")  
>>> checkpoint_location = "tmp/orders_checkpoint"  
>>>  
>>> keyspace = "streaming_1004"  
>>>  
>>>
```

```
>>> kafka_brokers = "bigdataanalytics2-worker-shdpt-v31-l-0:6667"
>>> kafka_topic = "orders_topic_json"
>>>
>>> raw_orders = spark.readStream. \
...     format("kafka"). \
...     option("kafka.bootstrap.servers", kafka_brokers). \
...     option("subscribe", kafka_topic). \
...     option("maxOffsetsPerTrigger", "20"). \
...     option("startingOffsets", "earliest"). \
...     load()
>>>
>>> schema = StructType() \
...     .add("order_id", StringType()) \
...     .add("customer_id", StringType()) \
...     .add("order_status", StringType()) \
...     .add("order_purchase_timestamp", StringType()) \
...     .add("order_approved_at", StringType()) \
...     .add("order_delivered_carrier_date", StringType()) \
...     .add("order_delivered_customer_date", StringType()) \
...     .add("order_estimated_delivery_date", StringType())
>>>
>>> parsed_orders = raw_orders \
...     .select(F.from_json(F.col("value").cast("String"), schema).alias("value"
...     ), "offset") \
...     .select("value.*")
>>>
>>>
```

```
>>> stream = console_output(parsed_orders, 10)
21/08/19 15:41:33 WARN shortcircuit.DomainSocketFactory: The short-circuit local
  reads feature cannot be used because libhadoop cannot be loaded.
-----
Batch: 0
-----
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|order_id|customer_id|order_status|
|order_purchase_timestamp|order_approved_at|order_delivered_carrier_date|order_
|delivered_customer_date|order_estimated_delivery_date|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|47770eb9100c2d0c44946d9cf07ec65d|41ce2a54c0b03bf3443c3d931a367089|delivered|
2018-08-08 08:38:49|2018-08-08 08:55:23|2018-08-08 13:50:00|2018-0
8-17 18:06:29|2018-09-04 00:00:00|
|ad21c59c0840e6cb83a9ceb5573f8159|8ab97904e6daea8866dbdbc4fb7aad2c|delivered|
2018-02-13 21:18:39|2018-02-13 22:20:29|2018-02-14 19:46:34|2018-0
2-16 18:17:02|2018-02-26 00:00:00|
|136cce7faa42fdb2cefd53fdc79a6098|ed0271e0b7da060a393796590e7b737a|invoiced|
2017-04-11 12:22:08|2017-04-13 13:25:17|
|2017-05-09 00:00:00|
```

```
6-16 15:20:55          |2018-07-18 00:00:00          |
|ecab90c9933c58908d3d6add7c6f5ae3|761df82feda9778854c6dafdaeb567e4|delivered |
2018-02-25 13:50:30    |2018-02-25 14:47:35|2018-02-26 22:28:50          |2018-0
3-27 23:29:14          |2018-04-13 00:00:00          |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
stream.stop()
>>>
```

Читаем из Cassandra, выводим схему и таблицу

```
>>> customer_names = spark.read \
...     .format("org.apache.spark.sql.cassandra") \
...     .options(keyspace=keyspace, table="customer_names") \
...     .load()
>>>
>>> customer_names.printSchema()
root
 |-- cid: string (nullable = true)
 |-- full_name: string (nullable = true)
>>> customer_names.show(truncate=False)
+-----+-----+
|cid                |full_name      |
+-----+-----+
|e4a32f8f3648818820a821cd577ccdba|Alison Martin  |
|75932cfd72b87bd6079cfl7786726807|Jill Lundrigan |
|d68585c54450af4bfc04cc6cccbcd607|James Mitchell |
|062328becf66582e5849b4f2b364b143|Sheikh Ahmed   |
|48f239e15744ed5e7ffbafc6bb6e882b|Kathy Denton   |
|5a58afc695ee03b9baca01d4afa52cec|Liam Walsh      |
|02a7a75320e787aa980907d0f647e508|Stephen Caller  |
|6772a0a230a2667dl6c3620f000e1348|Steven Marlow   |
|83da8aec5d2e8b2847e6ca45bea5588f|Gary Kang       |
|4c9c7c2b6de6ee2568681b5599bb7495|Dixon Coker     |
|65c295c5ac6110fc4214f32ab7df512e|Jonathan Craddock|
|93f0c8eaf1b7alc951d234d56b232e35|Gregory Watts   |
+-----+-----+
```

Связываем таблицы и выводим

```
>>> orders_by_names = parsed_orders\
...     .join(customer_names, parsed_orders.customer_id == customer_names.cid, "
inner")\
...     .select("order_id", customer_names.cid, customer_names.full_name)\
...     .withColumnRenamed("order_id", "oid")
>>>
>>> stream = console_output(orders_by_names, 10)
Batch: 0
+-----+-----+
+
|oid                |cid                |full_name      |
+-----+-----+
+
|76c6e866289321a7c93b82b54852dc33|f54a9f0e6b351c431402b8461ea51999|George Rae      |
|5ff96c15d0b717ac6ad1f3d77225a350|19402a48fe860416adf93348aba37740|Wanda Woods     |
|e481f51cbdc54678b7cc49136f2d6af7|9ef432eb6251297304e76186b10a928d|Kevin Maguire    |
+-----+-----+
```

Делаем выборку

```
>>> names_df = spark.sql("""select '20b5aae6a3e31111009f9a7ecc31a232' as cid, 'Ann Peterson 2' as full_name""")
>>> names_df.show()
+-----+-----+
|          cid|    full_name|
+-----+-----+
|20b5aae6a3e311110...|Ann Peterson 2|
+-----+-----+

>>>
```

Выводим таблицу

```
>>> all_names_df = spark.read \
...     .format("org.apache.spark.sql.cassandra") \
...     .options(keyspace=keyspace, table="customer_names") \
...     .load()
>>>
>>> all_names_df.show()
+-----+-----+
|          cid|    full_name|
+-----+-----+
|e4a32f8f364881882...|    Alison Martin|
|75932cfd72b87bd60...|    Jill Lundrigan|
|d68585c54450af4bf...|    James Mitchell|
|062328becf66582e5...|    Sheikh Ahmed|
|48f239e15744ed5e7...|    Kathy Denton|
|5a58afc695ee03b9b...|    Liam Walsh|
|02a7a75320e787aa9...|    Stephen Caller|
|6772a0a230a2667dl...|    Steven Marlow|
|83da8aec5d2e8b284...|    Gary Kang|
|4c9c7c2b6de6ee256...|    Dixon Coker|
|65c295c5ac6110fc4...|Jonathan Craddock|
|93f0c8eaf1b7alc95...|    Gregory Watts|
|f54a9f0e6b351c43l...|    George Rae|
|241e78de29b3090cf...|    Russell Carr|
|7c5f8c2e42af586aa...|    Will Harvey|
|20b5aae6a3e311110...|    Ann Peterson 2|
+-----+-----+
```

Фильтруем по полю cid и подсчитываем количество

```
>>> c_name_df = all_names_df.filter(F.col("cid") == "20b5aae6a3e31111009f9a7ecc31a232")
>>> c_name_df.show()
+-----+-----+
|          cid|    full_name|
+-----+-----+
|20b5aae6a3e311110...|Ann Peterson 2|
+-----+-----+

>>> c_name_df.count()
1
>>>
```

Фильтруем по полю full\_name и подсчитываем количество

```
>>> jane_df = all_names_df.filter(F.col("full_name") == "Ann Peterson")
>>> jane_df.show()
+---+-----+
|cid|full_name|
+---+-----+
+---+-----+

>>> jane_df.count()
0
>>>
```

Просматриваем план запроса

```
>>> jane_df.explain(True)
== Parsed Logical Plan ==
'Filter ('full_name = Ann Peterson)
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Analyzed Logical Plan ==
cid: string, full_name: string
Filter (full_name#496 = Ann Peterson)
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Optimized Logical Plan ==
Filter (isnotnull(full_name#496) && (full_name#496 = Ann Peterson))
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Physical Plan ==
*(1) Filter (isnotnull(full_name#496) && (full_name#496 = Ann Peterson))
+- *(1) Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b [cid#495,full_name#496] PushedFilters: [IsNotNull(full_name), EqualTo(full_name,Ann Peterson)], ReadSchema: struct<cid:string,full_name:string>
>>>
```

Задаем фильтр для набора данных и просматриваем план запроса

```
>>> between_select = all_names_df.filter(F.col("cid").between('20b5aae6a3e31111009f9a7ecc31a232', 'b89010d4a6acaa06d4ef89043869838e'))
>>> between_select.explain(True)
== Parsed Logical Plan ==
'Filter (('cid >= 20b5aae6a3e31111009f9a7ecc31a232) && ('cid <= b89010d4a6acaa06d4ef89043869838e))
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Analyzed Logical Plan ==
cid: string, full_name: string
Filter ((cid#495 >= 20b5aae6a3e31111009f9a7ecc31a232) && (cid#495 <= b89010d4a6acaa06d4ef89043869838e))
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b
```

Выводим набор данных и подсчитываем количество

```
>>> between_select.show()
+-----+
|          cid|          full_name|
+-----+
|75932cfd72b87bd60...|    Jill Lundrigan|
|48f239e15744ed5e7...|    Kathy Denton|
|5a58afc695ee03b9b...|    Liam Walsh|
|6772a0a230a2667d1...|    Steven Marlow|
|83da8aec5d2e8b284...|    Gary Kang|
|4c9c7c2b6de6ee256...|    Dixon Coker|
|65c295c5ac6110fc4...|Jonathan Craddock|
|93f0c8eaf1b7alc95...|    Gregory Watts|
|241e78de29b3090cf...|    Russell Carr|
|7c5f8c2e42af586aa...|    Will Harvey|
|20b5aae6a3e311110...|    Ann Peterson 2|
|6c347ef65dd574fb9...| Archie MacGregor|
|7e016f9ea27527978...|    Maurice Savage|
|ae2164e850f39dce4...|    Mohammed Hoque|
|4ad5a269a2d59d2c8...|    Amy Gillespie|
|90d2e6f72916e7282...|    Edward Lees|
|72600f002ba1550d6...|    Gordon Cutts|
|2fdffca8dcdf01547...|    John Collier|
|a00009bf8489ae779...|    Simon Broomhead|
|254e327149dc394a4...|    Dawn Connell|
+-----+
only showing top 20 rows

>>> between_select.count()
71
>>>
```

Формируем следующий набор данных и просматриваем план запроса

```
>>> in_select = all_names_df.filter(F.col("cid").isin('20b5aae6a3e31111009f9a7ecc31a232', 'b89010d4a6acaa06d4ef89043869838e'))
>>> in_select.explain(True)
== Parsed Logical Plan ==
'Filter 'cid IN (20b5aae6a3e31111009f9a7ecc31a232,b89010d4a6acaa06d4ef89043869838e)
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Analyzed Logical Plan ==
cid: string, full_name: string
Filter cid#495 IN (20b5aae6a3e31111009f9a7ecc31a232,b89010d4a6acaa06d4ef89043869838e)
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Optimized Logical Plan ==
Filter cid#495 IN (20b5aae6a3e31111009f9a7ecc31a232,b89010d4a6acaa06d4ef89043869838e)
+- Relation[cid#495,full_name#496] org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b

== Physical Plan ==
*(1) Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@6d9185b [cid#495,full_name#496] PushedFilters: [*In(cid, [20b5aae6a3e31111009f9a7ecc31a232,b89010d4a6acaa06d4ef89043869838e])], ReadSchema: struct<cid:string,full_name:string>
>>>
```



Выводим набор данных и подсчитываем количество записей

```
>>> in_select.show()
+-----+-----+
|          cid|    full_name|
+-----+-----+
|20b5aae6a3e311110...|Ann Peterson 2|
|b89010d4a6acaa06d...|  Stephen Wood|
+-----+-----+

>>> in_select.count()
21/08/19 16:08:33 WARN core.RequestHandler: Query '[2 bound values] SELECT count
(*) FROM "streaming_1004"."customer_names" WHERE "cid" IN (?, ?)  ALLOW FILTERI
NG;' generated server side warning(s): Aggregation query used on multiple partit
ion keys (IN restriction)
2
>>>
```

Задаем другой формат вывода, делаем выборку и выводим план запроса

```
>>> cass_big_df = spark.read \
...     .format("org.apache.spark.sql.cassandra") \
...     .options(table="users_many", keyspace="keyspace1") \
...     .load()
>>>
>>> between_select = cass_big_df.filter(F.col("user_id").between(4164237664, 416
4237664+10) )
>>> between_select.explain(True)
== Parsed Logical Plan ==
'Filter (('user_id >= 4164237664) && ('user_id <= 4164237674))
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSou
ceRelation@66e44222

== Analyzed Logical Plan ==
user_id: string, gender: string
Filter ((cast(user_id#576 as bigint) >= 4164237664) && (cast(user_id#576 as bigi
nt) <= 4164237674))
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSou
ceRelation@66e44222

== Optimized Logical Plan ==
Filter ((isnotnull(user_id#576) && (cast(user_id#576 as bigint) >= 4164237664))
&& (cast(user_id#576 as bigint) <= 4164237674))
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSou
ceRelation@66e44222

== Physical Plan ==
*(1) Filter (((cast(user_id#576 as bigint) >= 4164237664) && (cast(user_id#576 a
s bigint) <= 4164237674)) && isnotnull(user_id#576))
+- *(1) Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@66e44222 [us
```

Выводим результат

```
>>> between_select.show()
+-----+-----+
|  user_id|gender|
+-----+-----+
|4164237671|    9|
|4164237669|    9|
|4164237665|    9|
|4164237672|    9|
|4164237674|    9|
|4164237670|    9|
|4164237673|    9|
|4164237666|    9|
|4164237664|    9|
|4164237667|    9|
|4164237668|    9|
+-----+-----+
```

Делаем выборку набора данных и просматриваем план запроса

```
>>> in_select = cass_big_df.filter(F.col("user_id").isin(4164237664, 4164237665))
>>> in_select.explain(True)
== Parsed Logical Plan ==
'Filter 'user_id IN (4164237664,4164237665)
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSourceRelation@66e44222

== Analyzed Logical Plan ==
user_id: string, gender: string
Filter cast(user_id#576 as string) IN (cast(4164237664 as string),cast(4164237665 as string))
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSourceRelation@66e44222

== Optimized Logical Plan ==
Filter user_id#576 IN (4164237664,4164237665)
+- Relation[user_id#576,gender#577] org.apache.spark.sql.cassandra.CassandraSourceRelation@66e44222

== Physical Plan ==
*(1) Scan org.apache.spark.sql.cassandra.CassandraSourceRelation@66e44222 [user_id#576,gender#577] PushedFilters: [*In(user_id, [4164237664,4164237665])], ReadSchema: struct<user_id:string,gender:string>
>>>
```

Выводим набор данных

```
>>> in_select.show()
+-----+-----+
| user_id|gender|
+-----+-----+
|4164237664|    9|
|4164237665|    9|
+-----+-----+
>>>
```